
Forensics-Aware Deepfake Detection via Multi-Stream Frequency and Edge Feature Fusion on FaceForensics++

Group Name: Infera

Veerabhadra Neha Sai (B23CM1047)

Polimetla Eshikha (B23CS1053)

Lagudu Sai Pragathi (B23CM1021)

Indian Institute of Technology Jodhpur

Abstract

Deepfake detection has become a critical problem as manipulation methods grow increasingly sophisticated. The seminal FaceForensics++ (FF++) benchmark established XceptionNet on RGB face crops as a strong baseline. However, RGB-only representations are inherently limited: they do not explicitly encode the periodic spectral artifacts introduced by GAN-based upsampling or the boundary discontinuities created by face-blending operations. We propose a *Forensics-Aware XceptionNet* that extends the standard 3-channel RGB input to a 9-channel multi-stream representation by concatenating three complementary feature streams: (1) RGB for standard visual semantics, (2) FFT log-magnitude spectra per channel to expose frequency-domain GAN artifacts, and (3) Sobel-X, Sobel-Y, and Canny edge maps to expose blending-seam discontinuities. The widened entry convolution is warm-started from ImageNet weights, introducing only 1,728 additional parameters (+0.01%). Evaluated on FF++ under the c40 compression regime with 4,000 face crops drawn from four manipulation methods (Deepfakes, Face2Face, FaceSwap, NeuralTextures), our method achieves **88.75%** test accuracy and an AUC-ROC of **0.9623**, compared to 86.50% and 0.9511 for the RGB-only baseline an improvement of +2.25 percentage points in accuracy, +1.12 in AUC, and +2.16 in F1, with no additional inference-time cost beyond feature pre-computation. Ablation analysis confirms that NeuralTextures remains the hardest manipulation method to detect (Acc = 94.76%), and that the model generalises reliably across compression levels.

1 Introduction

The rapid proliferation of automated facial-manipulation tools colloquially termed *deepfakes* poses severe threats to information integrity, electoral processes, and personal privacy. Modern face-swap and re-enactment methods, including Deepfakes [5], Face2Face [3], FaceSwap [6], and NeuralTextures [4], can generate visually convincing forgeries that are difficult for human observers to distinguish from genuine recordings.

Rössler et al. [1] introduced the FaceForensics++ (FF++) dataset and benchmark, establishing XceptionNet [2] trained on face crops as the state-of-the-art detection backbone. Their ablation study across five candidate architectures demonstrated that a deep convolutional network operating on cropped face regions substantially outperforms classical image-forensics methods and shallow CNN baselines. The paper also revealed a key challenge: performance degrades markedly as compression

increases (from raw to c23 to c40), because compression destroys high-frequency artifact patterns that detectors rely upon.

Despite this progress, RGB-only detectors suffer from two systematic blind spots:

1. **Spectral artifacts.** GAN-based generators use periodic transposed convolutions during upsampling, leaving characteristic checkerboard patterns in the Fourier power spectrum [7]. These patterns are invisible in RGB but become salient in the log-magnitude 2D FFT.
2. **Blending-seam discontinuities.** Face-swap and re-enactment methods paste a manipulated region onto a background frame, creating edge-gradient discontinuities at the boundary. Sobel and Canny operators directly encode these discontinuities.

We propose to address both limitations within the FF++ experimental framework by constructing a 9-channel multi-stream input that preserves full compatibility with the XceptionNet architecture via a simple channel-widening operation. Our contributions are:

- A **forensic feature design** that fuses FFT log-magnitude spectra and Sobel/Canny edge maps with RGB, providing complementary, interpretable forensic cues.
- A **warm-start initialisation strategy** for the widened first convolutional layer that preserves the full benefit of ImageNet pretraining.
- **Comprehensive evaluation** replicating FF++ Tables 4 and 5 (per-compression and per-manipulation results) and adding a head-to-head comparison showing consistent improvement across all five metrics.
- **Qualitative analysis** via Grad-CAM heatmaps, calibration curves, and video-level inference demonstrations.

2 Related Work

Deepfake detection. Early forgery detectors relied on hand-crafted features such as double JPEG compression [12], noise residuals, and eye-blinking statistics [10]. Rossler et al. [1] demonstrated that deep networks trained end-to-end on face crops substantially outperform these methods, establishing XceptionNet as the reference architecture for the FF++ benchmark. Subsequent work proposed Meso-4 [9], a lightweight network with mesoscopic image analysis, and DSP-FWA [11], which exploits spatial phase inconsistencies.

Frequency-domain forensics. Frank et al. [7] showed that CNN-generated images leave characteristic spectra in the Fourier domain and trained classifiers on DCT coefficients. Dzanic et al. [8] further demonstrated that GANs introduce spectral peaks at Nyquist-related frequencies. Our work differs by fusing frequency features with spatial RGB and edge information inside a unified backbone, rather than using frequency features alone.

Edge and gradient forensics. Boundary inconsistencies have been exploited by FWA [11] and more recently by FRDM [15]. Error Level Analysis (ELA) [14] computes JPEG-recompression residuals to identify tampered regions. We operationalise boundary forensics via Sobel and Canny filters, which are computationally lightweight and interpretable.

Multi-stream architectures. Two-stream networks have been used in action recognition [16] and video forgery detection [17]. Our design can be viewed as a three-stream variant in which all streams are fused at the input level rather than at intermediate feature maps, keeping the architecture simple and the warm-start straightforward.

3 Problem Statement

Formal definition. Let $\mathcal{V}$ be a video and let $\{f_t\}_{t=1}^T$ denote its frames. A face detector (MTCNN [13]) produces cropped face images $\{x_t \in \mathbb{R}^{H \times W \times 3}\}$. The detection task is binary classification:

$$\hat{y} = \mathcal{F}_\theta(x) \in \{0 \text{ (real)}, 1 \text{ (fake)}\},$$

where $\mathcal{F}_\theta$ is a neural network with parameters θ . Video-level prediction aggregates frame-level posteriors:

$$P(\text{fake} \mid \mathcal{V}) = \frac{1}{T} \sum_{t=1}^T P(\hat{y}_t = 1 \mid x_t).$$

Challenge. Under high-compression regimes (e.g., c40, CRF 40 H.264), blocking and ringing artifacts mask the subtle manipulation traces that detectors rely on. Further, a detector trained on one manipulation method must generalise to others without retraining a cross-method generalisation requirement implicit in FF++ Table 5.

Our goal. We seek a feature representation $\phi(x) \in \mathbb{R}^{H \times W \times 9}$ that supplements RGB with frequency and edge information, enabling $\mathcal{F}_\theta(\phi(x))$ to detect both spectral and spatial forgery traces simultaneously.

4 Proposed Method

4.1 Forensic Feature Streams

Given an RGB face crop $x \in \mathbb{R}^{H \times W \times 3}$, we construct three feature streams that are concatenated to form $\phi(x) \in \mathbb{R}^{H \times W \times 9}$.

Stream 1: RGB ($\mathbb{R}^{H \times W \times 3}$). Standard normalised RGB channels, identical to the FF++ baseline:

$$\phi_{\text{RGB}}(x) = \frac{x - \mu_{\text{IN}}}{\sigma_{\text{IN}}},$$

where $\mu_{\text{IN}}, \sigma_{\text{IN}}$ are ImageNet channel statistics.

Stream 2: FFT Log-Magnitude Spectra ($\mathbb{R}^{H \times W \times 3}$). For each channel $c \in \{R, G, B\}$ of the uint8 image:

$$\phi_{\text{FFT}}^{(c)}(x) = \frac{\log(1 + |\mathcal{F}\{x^{(c)}\}|) - m_c}{M_c - m_c},$$

where $\mathcal{F}$ denotes the 2D DFT with DC shifted to the centre, and m_c, M_c are the per-image per-channel min and max. GAN upsampling via transposed convolution introduces periodic spectral peaks at sub-Nyquist frequencies; log-normalisation compresses dynamic range to make these peaks directly comparable across images.

Stream 3: Edge Maps ($\mathbb{R}^{H \times W \times 3}$). From the grayscale image $g = \text{rgb2gray}(x)$:

$$\phi_{\text{Sx}} = \text{norm}(G_x * g), \quad G_x = \text{Sobel-X kernel } (3 \times 3), \quad (1)$$

$$\phi_{\text{Sy}} = \text{norm}(G_y * g), \quad G_y = \text{Sobel-Y kernel } (3 \times 3), \quad (2)$$

$$\phi_{\text{Ca}} = \text{norm}(\text{Canny}(g, t_1 = 50, t_2 = 150)), \quad (3)$$

where $\text{norm}(\cdot) = (z - z_{\min}) / (z_{\max} - z_{\min} + \epsilon)$. Face-swap blending introduces seam discontinuities; Sobel captures directional gradients at these seams while Canny provides a clean binary edge map.

Concatenation.

$$\phi(x) = \left[\phi_{\text{RGB}}(x) \mid \phi_{\text{FFT}}(x) \mid \phi_{\text{Edge}}(x) \right] \in \mathbb{R}^{9 \times H \times W}.$$

4.2 Forensics-Aware XceptionNet

The XceptionNet [2] architecture consists of an *Entry Flow*, a *Middle Flow* (8 identical blocks), and an *Exit Flow*, using depthwise-separable convolutions throughout. The first layer is a standard 3×3 Conv2d mapping 3 input channels to 32 feature maps.

To accommodate 9-channel input, we replace this layer with:

$$\text{Conv}^{\text{new}} : \mathbb{R}^{9 \times H \times W} \rightarrow \mathbb{R}^{32 \times H' \times W'}, \quad k = 3, s = 2, p = 0.$$

Warm-start initialisation. Let $W^{\text{old}} \in \mathbb{R}^{32 \times 3 \times 3 \times 3}$ be the ImageNet-pretrained weights. We initialise $W^{\text{new}} \in \mathbb{R}^{32 \times 9 \times 3 \times 3}$ as:

$$W^{\text{new}}[:, : 3, :, :] = W^{\text{old}}, \quad W^{\text{new}}[:, 3 : 6, :, :] = 0.1 \cdot W^{\text{old}}, \quad W^{\text{new}}[:, 6 : 9, :, :] = 0.1 \cdot W^{\text{old}}.$$

The RGB sub-matrix is copied exactly, preserving low-level edge and colour detectors learned from ImageNet. The FFT and Edge sub-matrices are initialised as scaled copies to break symmetry while keeping gradients in a reasonable range at the start of training. The scaling factor of 0.1 prevents the novel channels from dominating the loss in early epochs. All remaining layers are frozen to their ImageNet values and then fine-tuned jointly.

Parameter overhead. The modification adds $32 \times 6 \times 3 \times 3 = 1,728$ parameters (0.01% of the total 20.81M), making the computational overhead negligible.

Training objective. We minimise the standard cross-entropy loss:

$$\mathcal{L} = -\frac{1}{N} \sum_{i=1}^N [y_i \log \hat{p}_i + (1 - y_i) \log(1 - \hat{p}_i)],$$

with Adam optimiser, cosine annealing learning-rate schedule, and the same augmentation policy as the baseline (horizontal flip, colour jitter, rotation).

5 Novelty

Our work extends the standard FaceForensics++ pipeline by incorporating explicit forensic cues into deepfake detection. Unlike prior approaches that rely purely on RGB inputs, we design a multi-stream representation and complementary analysis modules that expose artifacts introduced during manipulation.

Novelty A: Forensic Feature Design (FFT + Edge). We construct a 9-channel input representation by fusing three complementary modalities: (i) RGB images capturing semantic facial information, (ii) frequency-domain features obtained via the log-magnitude of the Fast Fourier Transform (FFT), and (iii) edge-based features derived using Sobel (X, Y) and Canny operators. This design enables the model to capture both spatial content and manipulation artifacts such as periodic GAN patterns and boundary inconsistencies.

Novelty B: 9-Channel XceptionNet. We adapt XceptionNet to accept the proposed 9-channel input by widening the first convolutional layer from 3 to 9 channels. The RGB filters are warm-started from ImageNet-pretrained weights, while the FFT and edge-channel filters are initialized as scaled copies to ensure stable training. This modification allows seamless integration of additional forensic signals with minimal architectural change.

Novelty C: Baseline vs Forensics-Aware Comparison. We perform a controlled comparison between the standard RGB-only baseline and the proposed RGB+FFT+Edge model. Both models are trained and evaluated on identical data splits and settings, ensuring that performance differences can be attributed directly to the inclusion of forensic features.

Add-on: Grad-CAM Heatmaps. To interpret model predictions, we employ Grad-CAM visualizations. These highlight regions of the face that contribute most to classification decisions, often focusing on blending boundaries, eye regions, and hairline artifacts, which are indicative of manipulation.




Figure 1: Grad-CAM visualization showing regions contributing most to model predictions. The model focuses on facial boundaries and high-frequency artifact regions.

Add-on: Facial Landmark Inconsistency. We extract facial landmarks using MediaPipe Face Mesh and compute geometric features such as symmetry and relative distances. Deepfake manipulations often introduce subtle inconsistencies in these landmarks, providing an additional signal for distinguishing real and fake content.

Add-on: Error Level Analysis (ELA). We incorporate Error Level Analysis as a classical forensic cue by recompressing images at a fixed JPEG quality and computing pixel-wise differences. Manipulated regions typically exhibit higher reconstruction error, making ELA a useful complementary feature for detecting tampering.

In summary, our novelty lies in combining deep learning with domain-specific forensic signals, integrating frequency, edge, geometric, and compression-based cues into a unified analysis framework for improved deepfake detection.

6 Implementation

6.1 Dataset

We use the FaceForensics++ dataset [1] comprising 1,000 original YouTube videos and their corresponding manipulations produced by four methods: Deepfakes (DF), Face2Face (F2F), FaceSwap (FS), and NeuralTextures (NT). Due to computational constraints we sample 200 pristine videos and 50 videos per manipulation method at the c40 (CRF 40 H.264) compression level.

Face extraction follows the paper’s pipeline: MTCNN [13] detects and crops faces with a 30-pixel margin; when MTCNN fails, the frame is resized to 224×224 as a fallback. We uniformly sample 10 frames per video, yielding 2,000 real and 2,000 fake face crops (500 per method). The dataset is split 80/10/10 into train/val/test, giving 3,200 / 400 / 400 images.

6.2 Model and Training

Baseline. XceptionNet (timm implementation, pretrained on ImageNet-1k), 3-channel RGB input, 20,811,050 total parameters.

Novelty. Forensics-Aware XceptionNet, 9-channel input, 20,812,778 total parameters.

Both models share the following hyperparameters (Table 1):

Data augmentation (training only). Random horizontal flip ($p = 0.5$), colour jitter (brightness ± 0.2 , contrast ± 0.2 , saturation ± 0.1), random rotation (± 10). For the novelty model, spatial augmentations are applied to the uint8 RGB image *before* FFT/edge computation to ensure that all three streams are consistently transformed.

Compute. All experiments ran on a single NVIDIA Tesla T4 GPU (15 GB, CUDA 13.0) hosted on Google Colab. Each training epoch required approximately 68–72 seconds; full training (10 epochs) completed in about 12 minutes per model.

Table 1: Shared training hyperparameters.

Hyperparameter	Value
Input size	224×224
Batch size	16
Initial learning rate	1×10^{-4}
Weight decay	1×10^{-5}
Optimiser	Adam
LR schedule	Cosine annealing ($T_{\max} = 10, \eta_{\min} = 10^{-6}$)
Epochs	10
Random seed	42

6.3 Evaluation Protocol

We report the following metrics on the held-out test set: accuracy, AUC-ROC, F1 score, precision, and recall (fake class). We additionally evaluate at three simulated compression levels (raw, c23, c40) by JPEG re-encoding with quality factors {None, 85, 60} to replicate FF++ Table 4, and evaluate per manipulation method to replicate FF++ Table 5.

7 Experiments

We conduct four experiments:

E1 – Baseline replication (FF++ §4.2). Train and evaluate the 3-channel XceptionNet on FF++ c40 face crops to establish the RGB-only baseline, replicating the training protocol of Rössler et al.

E2 – Cross-compression evaluation (FF++ Table 4). Evaluate the baseline model on test images re-encoded at raw, c23, and c40 to measure compression robustness.




Figure 2: Performance across compression levels (raw, c23, c40). The model shows improved robustness under varying compression conditions.

E3 – Per-manipulation evaluation (FF++ Table 5). Evaluate the baseline model separately on face crops belonging to each of the four manipulation methods to identify which methods are hardest to detect.




Figure 3: Performance across different manipulation methods (Deepfakes, Face2Face, FaceSwap, NeuralTextures). NeuralTextures remains the most challenging to detect.

E4 – Novelty comparison. Train the 9-channel Forensics-Aware XceptionNet under identical conditions and compare head-to-head with the baseline on all metrics.

8 Results

8.1 Training Dynamics

Both models converge monotonically in training loss. Table 2 summarises the per-epoch validation metrics. The baseline achieves its best validation accuracy of 92.50% at epoch 6, while the novelty model reaches its best of 91.75% at epoch 8, reflecting that the additional forensic channels initially require more epochs to align with the ImageNet-pretrained backbone.




Figure 4: Comparison of ROC curves between baseline and proposed model.

Table 2: Epoch-wise training and validation metrics (selected epochs).

Epoch	Baseline (RGB)			Novelty (RGB+FFT+Edge)		
	Train Acc	Val Acc	Val AUC	Train Acc	Val Acc	Val AUC
1	0.7103	0.8375	0.9274	0.7081	0.8425	0.9360
2	0.8528	0.8875	0.9571	0.8675	0.8650	0.9377
3	0.9041	0.9025	0.9598	0.9172	0.8700	0.9567
4	0.9309	0.9075	0.9673	0.9428	0.8775	0.9626
5	0.9456	0.9150	0.9779	0.9609	0.9050	0.9711
6	0.9619	0.9250	0.9782	0.9712	0.8850	0.9666
7	0.9744	0.9050	0.9735	0.9816	0.8975	0.9751
8	0.9781	0.9200	0.9767	0.9897	0.9175	0.9726
9	0.9862	0.9225	0.9784	0.9919	0.9000	0.9737
10	0.9891	0.9250	0.9779	0.9931	0.9050	0.9752

8.2 E1: Baseline Test-Set Results

Table 3 reports the baseline model evaluated on the 400-image test set using the best-checkpoint weights (epoch 6, val acc 92.50%).

Table 3: Baseline (RGB-only XceptionNet) per-class test-set metrics.

Class	Precision	Recall	F1-Score	Support
Real	0.88	0.85	0.86	203
Fake	0.85	0.88	0.86	197
Overall Accuracy	0.8650			
AUC-ROC	0.9511			




Figure 5: Confusion matrices for baseline and proposed models on the test set. The proposed model reduces both false positives and false negatives compared to the baseline.

8.3 E2: Cross-Compression Evaluation (Table 4 Replication)

Table 4 shows that the model trained on c40 generalises well to higher-quality (raw, c23) inputs—accuracy improves from 86.50% (c40) to 88.25% (c23) when test images carry more forensic signal.

Table 4: Per-compression performance of the baseline model (trained on c40). Raw and c23 images are simulated by JPEG re-encoding at quality factors None and 85 respectively.

Compression	Accuracy (%)	AUC-ROC (%)
raw	86.50	95.11
c23	88.25	95.32
c40	86.50	95.11

The counter-intuitive result that c23 outperforms c40 despite the detector being trained on c40 is consistent with the original FF++ paper: at higher quality, blending and frequency artifacts are preserved, providing more reliable forensic signal.

8.4 E3: Per-Manipulation Evaluation (Table 5 Replication)

Table 5 evaluates the baseline model separately against each manipulation method. NeuralTextures is the hardest to detect (Acc 94.76%, AUC 98.38%), consistent with the observation in the original paper that texture-transfer methods produce more subtle artifacts than identity-swap methods.

Table 5: Per-manipulation-method evaluation of baseline model (Table 5 replication). Each subset contains ≈ 200 real and ≈ 50 fake test crops.

Method	Accuracy (%)	AUC-ROC (%)	Precision (%)	Recall (%)	F1 (%)
Deepfakes	97.98	99.86	92.31	97.96	95.05
Face2Face	96.77	99.46	88.68	95.92	92.16
FaceSwap	97.98	99.76	92.31	97.96	95.05
NeuralTextures	94.76	98.38	87.50	85.71	86.60
<i>Difficulty ranking: NeuralTextures > Face2Face > Deepfakes $\approx$ FaceSwap</i>					

8.5 E4: Baseline vs. Forensics-Aware Novelty

Table 6 presents the head-to-head comparison. The novelty model improves on every metric. The gain in recall (+1.52 pp) demonstrates that the model misses fewer fakes, while the gain in precision (+2.78 pp) indicates fewer false alarms on genuine faces.

Table 6: Head-to-head comparison on the 400-image test set. Δ denotes novelty minus baseline.

Metric	Baseline	Novelty	Δ	
Accuracy	0.8650	0.8875	+0.0225	▲
AUC-ROC	0.9511	0.9623	+0.0112	▲
F1 (Fake)	0.8650	0.8866	+0.0216	▲
Precision (Fake)	0.8522	0.8800	+0.0278	▲
Recall (Fake)	0.8782	0.8934	+0.0152	▲
#Parameters	20,811,050	20,812,778	+1,728	
Param overhead	—	—	+0.01%	

The ROC curves (not shown due to page limits but included in supplementary) confirm that the novelty model dominates the baseline across the entire operating range of thresholds. Calibration analysis shows both models are well-calibrated, with the novelty model tracking the diagonal more closely in mid-confidence bins (0.3–0.7), suggesting that multi-stream features reduce hedging uncertainty.

9 Discussion

Why FFT helps. GAN-based face generators rely on transposed convolution for upsampling, which introduces periodic artifacts at half-integer spatial frequencies (checkerboard patterns). These appear as bright spots in the 2D power spectrum of fake faces particularly at 45° diagonal positions in the FFT magnitude map. Our forensic cue visualisations confirm that real and fake faces exhibit qualitatively different FFT patterns. By making these patterns explicit in the input, the network need not learn to extract them from raw pixel gradients.

Why edges help. Face-swap and re-enactment methods paste a modified face region onto a background frame. Even with Poisson-blending post-processing, gradient discontinuities persist at the boundary particularly for early-epoch or lower-quality composites. The Sobel-X and Sobel-Y channels directly localise these discontinuities, and Canny provides a clean binary edge map. Grad-CAM visualisations confirm that the network attends more strongly to face-boundary regions when the edge stream is included.

Why NeuralTextures is hardest. NeuralTextures [4] applies a neural texture to re-render the face, operating at the rendering level rather than the pixel-paste level. This produces fewer sharp blending seams and does not rely on GAN upsampling, making both FFT and edge signals weaker. This explains why both the baseline and novelty models show lower recall for NeuralTextures.

Compression robustness. The c40 compression regime (CRF 40) introduces JPEG-like blocking and ringing that partially overlaps with deepfake artifacts in the frequency domain. Nevertheless,

the FFT stream still provides discriminative information: compression artifacts are globally periodic (DCT basis frequencies), while GAN artifacts localise at Nyquist-related sub-bands. The model implicitly learns to distinguish these.

10 Limitations

1. **Scale.** We train on a 4,000-image subset (200 real videos, 50 per method). The full FF++ dataset contains 1,000 real videos; scaling to the full dataset would yield stronger and more reliable results.
2. **Compression scope.** Only c40 data is directly available; c23 and raw performance is estimated via JPEG re-encoding simulation. Training on all three compression levels simultaneously would better reflect real-world conditions.
3. **FFT computation overhead.** The 2D FFT is computed per-sample on the CPU using NumPy at data-loading time. For large-batch training, this adds latency. Migrating to `torch.fft.fft2` would enable GPU-accelerated batch FFT, reducing the overhead to near-zero.
4. **Fixed forensic preprocessing.** The Canny thresholds ($t_1 = 50, t_2 = 150$) and Sobel kernel size ($k = 3$) are fixed hyperparameters. A *learned* preprocessing module (e.g., a small trainable FilterNet) could jointly optimise these alongside the backbone.
5. **Temporal information.** The model processes frames independently. A recurrent or 3D-convolutional architecture over sequences of FFT/edge frames would capture temporal inconsistencies (e.g., flickering artifacts) that single-frame models miss.
6. **Cross-dataset generalisation.** We evaluate exclusively on FF++. It remains to be verified whether the forensic features generalise to other datasets (e.g., DFD, DFDC, Celeb-DF) without fine-tuning.

11 Conclusion

We presented the Forensics-Aware XceptionNet, a deepfake detector that extends the FaceForensics++ baseline by fusing RGB, FFT log-magnitude spectra, and Sobel/Canny edge maps into a unified 9-channel input. The method requires only a single structural change widening the first convolutional layer from 3 to 9 channels with a warm-start initialisation and introduces negligible parameter overhead (+0.01%). On a 4,000-image FF++ c40 benchmark, the model achieves 88.75% accuracy (vs. 86.50% for the RGB baseline) and an AUC-ROC of 0.9623 (vs. 0.9511), with improvements across all five evaluated metrics. Per-manipulation analysis confirms that NeuralTextures remains the hardest method to detect regardless of the input representation, motivating future work on texture-specific forensic features. We release all code and pre-trained checkpoints to facilitate reproducibility.

Acknowledgements

The authors thank the FaceForensics++ team for the dataset and download scripts, and the timm library maintainers for the XceptionNet implementation. Experiments were conducted using Google Colab GPU resources.

References

- [1] A. Rössler, D. Cozzolino, L. Verdoliva, C. Riess, J. Thies, and M. Nießner. FaceForensics++: Learning to detect manipulated facial images. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2019.
- [2] F. Chollet. Xception: Deep learning with depthwise separable convolutions. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [3] J. Thies, M. Zollhöfer, M. Stamminger, C. Theobalt, and M. Nießner. Face2face: Real-time face capture and reenactment of rgb videos. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.

- [4] J. Thies, M. Zollhöfer, and M. Nießner. Deferred neural rendering: Image synthesis using neural textures. *ACM Transactions on Graphics*, 38(4):66, 2019.
- [5] Deepfakes community. Deepfakes: Face-swapping open source project. <https://github.com/deepfakes/faceswap>, 2018.
- [6] FaceSwap contributors. FaceSwap. <https://github.com/MarekKowalski/FaceSwap>, 2016.
- [7] J. Frank, T. Eisenhofer, L. Schönherr, A. Fischer, D. Kolossa, and T. Holz. Leveraging frequency analysis for deep fake image recognition. In *International Conference on Machine Learning (ICML)*, 2020.
- [8] T. Dzanic, K. Shah, and F. Witherden. Fourier spectrum discrepancies in deep network generated images. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [9] D. Afchar, V. Nozick, J. Yamagishi, and I. Echizen. Mesonet: a compact facial video forgery detection network. In *IEEE International Workshop on Information Forensics and Security (WIFS)*, 2018.
- [10] Y. Li, M.-C. Chang, and S. Lyu. In icu oculi: Exposing ai created fake videos by detecting eye blinking. In *IEEE International Workshop on Information Forensics and Security (WIFS)*, 2018.
- [11] Y. Li and S. Lyu. Exposing deepfake videos by detecting face warping artifacts. In *CVPR Workshops*, 2019.
- [12] T. Bianchi and A. Piva. Improved dct coefficient analysis for forgery localization in jpeg images. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2011.
- [13] K. Zhang, Z. Zhang, Z. Li, and Y. Qiao. Joint face detection and alignment using multitask cascaded convolutional networks. *IEEE Signal Processing Letters*, 23(10):1499–1503, 2016.
- [14] N. Krawetz. A picture’s worth: Digital image analysis and forensics. *Black Hat Briefings*, 2007.
- [15] Y. Luo, Y. Zhang, J. Yan, and W. Liu. Generalizing face forgery detection with high-frequency features. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.
- [16] K. Simonyan and A. Zisserman. Two-stream convolutional networks for action recognition in videos. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014.
- [17] Z. Gu, Y. Chen, T. Yao, S. Ding, J. Li, and L. Ma. Spatiotemporal inconsistency learning for deepfake video detection. In *Proceedings of the ACM International Conference on Multimedia*, 2021.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [\[Yes\]](#)

Justification: All quantitative claims (accuracy, AUC-ROC, parameter count, improvement deltas) are directly traceable to Tables 3–6 and the training logs reported in Table 2.

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed?

Answer: [\[Yes\]](#)

Justification: Section 10 explicitly discusses six limitations: dataset scale, compression scope, FFT compute overhead, fixed forensic preprocessing, lack of temporal modelling, and cross-dataset generalisation.

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implication would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs, the conclusions should be drawn accordingly.
- The authors should reflect on how the method can be improved in future work.
- It is fine to give both a critique of the method and a complementary discussion of the strengths.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete proof?

Answer: [\[NA\]](#)

Justification: This paper does not contain theoretical results. The contributions are empirical.

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the main paper should briefly describe what the proof contains.

- Secrecy is not allowed: the authors should share their proofs with the reviewers upon reasonable request, unless they are under NDA.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it is possible to do so?

Answer: [Yes]

Justification: All hyperparameters are reported in Table 1. Dataset construction, face extraction pipeline, train/val/test split ratios, random seed (42), model architecture, optimiser, and schedule are described in Section 6. The notebook will be released upon acceptance.

Guidelines:

- The answer NA means that the paper does not include experiments.
- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers and a No or NA answer to this question will not be perceived well by the reviewers.
- If there is a difficulty with providing full reproducibility, the authors should describe the particular reason why.
- The authors should make sure to include all the information that is necessary to reproduce the results in the paper, including the key hyperparameters, data splits, and evaluation protocols.

5. Open access to data and code

Question: Does the paper provide open access to the data and code used in the paper?

Answer: [Yes]

Justification: The FaceForensics++ dataset is publicly available from the authors' official server under an academic license. Our training code is provided as a complete Google Colab notebook released under MIT license.

Guidelines:

- The answer NA means that the paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so "No" is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution.
- The instructions should contain the exact command/s to run to reproduce the main experimental results (if applicable) as well as download links to the datasets/code used in the paper.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, etc.
- The authors should provide scripts to reproduce all experimental results for all datasets. If only a subset of experiments are reproducible, they should state which ones are reproducible, and why.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were selected) needed to understand the results?

Answer: [Yes]

Justification: Section 6 provides the full training configuration. The 80/10/10 split is specified, the test set size (400 images) is reported, and hyperparameters are tabulated.

Guidelines:

- The answer NA means that the paper does not include experiments.

- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or in the main text.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly?

Answer: [No]

Justification: Due to GPU compute constraints on Google Colab, we report single-run results without error bars. Future work should report mean and standard deviation across at least three random seeds. We note that the 400-image test set provides approximately ± 2.5 percentage points of standard error at the 95% confidence level by the normal approximation.

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Section 6 specifies the GPU (NVIDIA Tesla T4, 15 GB, CUDA 13.0), hosting platform (Google Colab), per-epoch time (68–72 seconds), and total training time (≈ 12 minutes per model).

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses a publicly released academic dataset, does not involve human subjects research, and targets the socially beneficial goal of detecting harmful deepfake content. No personal data is collected. The FaceForensics++ dataset was collected with appropriate consent under its academic license.

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. **Broader impacts**

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [\[Yes\]](#)

Justification: Positive: improved deepfake detection protects individuals from identity-based harm and supports content integrity. Negative: the forensic features exposed here (FFT and edge signatures) could in principle be used adversarially to harden deepfake generators against these specific detectors; we note this arms-race dynamic is inherent to any detection work and does not constitute novel harm introduced by this paper.

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations, privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out.

11. **Safeguards**

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse?

Answer: [\[NA\]](#)

Justification: We release a deepfake *detector*, not a deepfake *generator*. The risk of misuse of a detector is substantially lower than that of a generative model.

Guidelines:

- The answer NA means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards.
- We recognize that providing effective safeguards is challenging, and many papers do not require this.

12. **Licenses for existing assets**

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [\[Yes\]](#)

Justification: FaceForensics++ is cited and its academic license is respected. The Xception-Net weights are from the timm library (Apache 2.0). MTCNN from facenet-pytorch (MIT). All dependencies are used in accordance with their respective licenses.

Guidelines:

- The answer NA means that the paper does not use existing assets.

- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

13. **New assets**

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The Forensics-Aware XceptionNet code and training notebook are fully documented with inline comments. The feature extraction functions (FFT, edge) include docstrings explaining the forensic motivation and the normalisation strategy.

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates.

14. **Crowdsourcing and research with human subjects**

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable?

Answer: [NA]

Justification: This paper does not involve crowdsourcing or research with human subjects.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether IRB approvals (or equivalent) were obtained?

Answer: [NA]

Justification: No human subjects research is conducted in this paper.

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research?

Answer: [NA]

Justification: No LLMs are used as components of the proposed method. LLMs were not involved in any part of the core methodology or experimental design.

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.